

---

# PERFORMANCE CROSSOVER POINT ANALYSIS BETWEEN HASH TABLE AND BLOOM FILTER UNDER STRICT MEMORY CONSTRAINTS

---

TECHNICAL REPORT

Group 2 — SE20B02<sup>1</sup>

<sup>1</sup>Data Structures & Algorithms (CSD201), FPT University, Summer 2026

July 3, 2026

## ABSTRACT

In the domain of membership query problems—determining whether a given element belongs to a predefined set—Hash Tables and Bloom Filters represent a classical space–accuracy trade-off. Hash Tables provide exact, deterministic lookups at the cost of significant per-element memory overhead, whereas Bloom Filters offer a compact probabilistic alternative that permits a controllable false positive rate (FPR). This paper presents a systematic empirical study that identifies the *Performance Crossover Point*: the dataset scale  $N$  at which a Hash Table, constrained by a hard memory budget  $M$ , begins to degrade catastrophically while a Bloom Filter continues to operate with stable throughput. We implement both data structures from scratch in Java—a Hash Table with Separate Chaining and a Bloom Filter with Kirsch–Mitzenmacher double hashing ( $k = 7$ , target FPR  $p = 1\%$ )—and benchmark them across dataset sizes from  $N = 1,000$  to  $N = 2,000,000$  under five JVM memory caps (16–256 MB). Our results demonstrate that the Hash Table consumes approximately 100 bytes per element versus 1.2 bytes for the Bloom Filter (an  $\approx 83\times$  compression ratio), and that JVM garbage collection pressure causes Hash Table throughput to collapse shortly before an `OutOfMemoryError`. The empirical FPR of the Bloom Filter remains tightly bounded within 0.84%–1.38%, closely matching the Mitzenmacher theoretical prediction. We provide concrete crossover thresholds for each memory tier and discuss practical implications for system design in CDN caching, database pre-filtering, and network deduplication.

**Keywords** Hash Table · Bloom Filter · Membership Query · Performance Crossover · Memory Constraint · False Positive Rate

## 1 Introduction

Membership testing—answering the question “Does element  $x$  belong to set  $S$ ?”—is among the most fundamental operations in computer science. It underpins a wide range of systems including database query engines, network routers, web caches, spell checkers, and distributed storage systems [Broder and Mitzenmacher, 2004, Chang et al., 2008]. Two classical data structures address this problem with starkly different resource profiles:

- (i) **Hash Table** (with collision resolution): Provides exact, deterministic  $O(1)$  amortized lookup with zero false positives, but requires storing every key in memory—resulting in substantial per-element overhead.
- (ii) **Bloom Filter** [Bloom, 1970]: A space-efficient probabilistic structure that encodes set membership in a compact bit array, reducing memory by orders of magnitude at the cost of a tunable false positive rate (FPR).

While the theoretical trade-offs between these structures are well understood [Mitzenmacher and Upfal, 2017, Cormen et al., 2022], practitioners often lack concrete empirical guidance on *when* the transition from one to the other becomes necessary. Specifically, given a fixed memory budget  $M$ , at what dataset scale  $N$  does the Hash Table begin to degrade so severely that the Bloom Filter—despite its probabilistic nature—becomes the superior choice?

## 1.1 Research Question

We formalize this inquiry as follows:

*Given a fixed memory budget  $M$ , at what dataset scale  $N$  does a Hash Table with Separate Chaining begin to exhibit catastrophic performance degradation (due to RAM exhaustion and garbage collection pressure), and does a Bloom Filter configured with  $k$  hash functions and  $m$  bits demonstrate superior throughput at that point, with what empirical trade-off in false positive rate  $p$ ?*

## 1.2 Contributions

This paper makes the following contributions:

1. We design and implement a controlled, memory-constrained benchmark framework in Java that isolates the crossover behavior between Hash Table and Bloom Filter under hard JVM heap limits.
2. We identify concrete *crossover points* for five memory tiers (16–256 MB), quantifying the dataset size  $N$  at which Hash Table throughput collapses.
3. We empirically validate the Mitzenmacher formula for optimal Bloom Filter parameters, confirming that the measured FPR (0.84%–1.38%) closely tracks the theoretical target of 1%.
4. We quantify the memory compression ratio at approximately  $83\times$ , demonstrating that a Bloom Filter storing  $N = 2,000,000$  elements consumes only 2.3 MB versus  $\approx 191$  MB for the Hash Table.

The remainder of this paper is organized as follows. Section 2 reviews the theoretical foundations of both data structures. Section 3 details our experimental design and implementation. Section 4 presents the empirical results with quantitative analysis. Section 5 interprets the crossover phenomenon and its practical implications. Section 6 concludes with recommendations and future work.

## 2 Background and Related Work

### 2.1 Hash Table with Separate Chaining

A hash table maps keys to array indices via a hash function  $h : U \rightarrow \{0, 1, \dots, m-1\}$ , where  $U$  is the universe of keys and  $m$  is the table capacity. When two distinct keys  $k_1 \neq k_2$  satisfy  $h(k_1) = h(k_2)$ , a *collision* occurs. Separate Chaining resolves collisions by maintaining a linked list at each bucket [Cormen et al., 2022].

Under the *simple uniform hashing assumption* (SUHA), the expected chain length equals the load factor  $\alpha = n/m$ , yielding an expected lookup time of  $\Theta(1 + \alpha)$ . When  $\alpha$  is kept bounded (e.g.,  $\alpha \leq 0.75$ ) via dynamic resizing, lookups remain  $O(1)$  amortized [Knuth, 1998].

**Memory Cost.** Each element requires storage for:

- The key string itself ( $\approx 40 + \ell$  bytes, where  $\ell$  is the string length),
- A linked-list node (object header + key reference + next pointer  $\approx 32$  bytes),
- A bucket reference in the table array (8 bytes on a 64-bit JVM).

For typical usernames ( $\ell \approx 12$ ), this totals approximately **100 bytes per element**.

### 2.2 Bloom Filter

A Bloom Filter [Bloom, 1970] is a probabilistic data structure that represents a set  $S$  of  $n$  elements using a bit array  $B[0 \dots m-1]$  and  $k$  independent hash functions  $h_1, h_2, \dots, h_k$ . To insert an element  $x$ , the filter sets bits  $B[h_i(x)] \leftarrow 1$  for all  $1 \leq i \leq k$ . To query membership, it checks whether all  $k$  bits are set: if any bit is 0, the element is *definitely absent*; if all bits are 1, the element is *probably present* (with a probability of error equal to the FPR).

**Optimal Parameters** (Mitzenmacher Formula [Mitzenmacher and Upfal, 2017]). Given  $n$  elements and a target false positive rate  $p$ , the optimal bit-array size and number of hash functions are:

$$m = \left\lceil -\frac{n \ln p}{(\ln 2)^2} \right\rceil \quad (1)$$

$$k = \left\lceil \frac{m}{n} \cdot \ln 2 \right\rceil \quad (2)$$

For  $p = 0.01$  (1%), these simplify to  $m \approx 9.585 \cdot n$  bits and  $k = 7$ . The resulting memory cost is:

$$\text{Memory} = \frac{m}{8} \approx \frac{9.585n}{8} \approx \mathbf{1.2 \text{ bytes per element}}$$

**No False Negatives.** A correctly implemented Bloom Filter guarantees zero false negatives: if  $x \in S$ , then  $\text{lookup}(x)$  always returns `true`.

### 2.3 Double Hashing Technique

Kirsch and Mitzenmacher [Kirsch and Mitzenmacher, 2006] proved that only *two* independent hash functions  $h_1$  and  $h_2$  are needed to simulate  $k$  hash functions without increasing the asymptotic FPR:

$$g_i(x) = (h_1(x) + i \cdot h_2(x)) \bmod m, \quad i = 0, 1, \dots, k-1 \quad (3)$$

This technique dramatically reduces computational cost—instead of evaluating  $k$  independent hash functions, only two hash computations are required per operation.

### 2.4 Related Work

Broder and Mitzenmacher [Broder and Mitzenmacher, 2004] provide a comprehensive survey of Bloom Filter applications in networking, including distributed caching, routing, and peer-to-peer systems. Fan et al. [Fan et al., 2014] propose the Cuckoo Filter as an alternative that supports deletion, though with higher per-element overhead. Chang et al. [Chang et al., 2008] describe the use of Bloom Filters in Google’s Bigtable to reduce unnecessary disk reads—a direct application of the crossover phenomenon we study.

Our work differs from prior studies by focusing specifically on the *transition boundary* under strict, quantified memory constraints, using the JVM’s `-Xmx` flag to simulate hard RAM budgets and observe the exact crossover dynamics.

## 3 Methodology

This section describes the implementation of both data structures, the benchmark framework, and the experimental setup.

### 3.1 Implementation Details

Both data structures are implemented from scratch in Java (no external libraries), adhering to the curriculum requirements of CSD201.

#### 3.1.1 Hash Table with Separate Chaining (HashTableSC)

Our Hash Table uses an array of singly-linked lists for collision resolution. Key design decisions include:

- **Prime-sized table:** The table capacity is always set to the next prime  $\geq n$ , improving hash distribution under modular compression [Cormen et al., 2022].
- **Load factor threshold:** Dynamic resizing triggers when  $\alpha > 0.75$ , doubling the capacity to the next prime.
- **Hash function:** Java’s `String.hashCode()` with bitmask (`& 0x7FFFFFFF`) to ensure non-negative indices.
- **Head insertion:** New nodes are prepended to the chain in  $O(1)$  time.

#### 3.1.2 Bloom Filter (BloomFilter)

The Bloom Filter uses a `boolean[]` array of  $m$  bits with  $k = 7$  hash functions generated via double hashing (Equation (3)):

- **Primary hash  $h_1$ :** Java’s `String.hashCode()` with bitmask.
- **Secondary hash  $h_2$ :** FNV-1a algorithm [Carter and Wegman, 1979] with forced odd output to ensure full-period coverage.
- **Parameters:** Computed dynamically using Equations (1) and (2) with target  $p = 0.01$ .

Algorithm 1 shows the lookup procedure.

**Algorithm 1** Bloom Filter Membership Query**Require:** Key  $x$ , bit array  $B[0..m-1]$ , hash functions  $h_1, h_2$ , count  $k$ **Ensure:** true (probably in set) or false (definitely not)

```

1: for  $i \leftarrow 0$  to  $k - 1$  do
2:    $pos \leftarrow (h_1(x) + i \cdot h_2(x)) \bmod m$ 
3:   if  $B[pos] = 0$  then
4:     return false ▷ Definitely not in set
5:   end if
6: end for
7: return true ▷ Probably in set

```

Table 1: Experimental Configuration

| Parameter               | Value                                           |
|-------------------------|-------------------------------------------------|
| Language & Runtime      | Java (OpenJDK), single-threaded                 |
| Memory Budgets $M$      | 16, 32, 64, 128, 256 MB (-Xmx)                  |
| Dataset Sizes $N$       | 1K, 5K, 10K, 25K, 50K, 100K, 250K, 500K, 1M, 2M |
| Data Format             | JSON arrays of unique username strings          |
| Hash Table Load Factor  | $\alpha_{\max} = 0.75$ (auto-resize)            |
| Hash Table Capacity     | Next prime $\geq N$                             |
| Bloom Filter Target FPR | $p = 0.01$ (1%)                                 |
| Bloom Filter Hash Count | $k = 7$ (computed from Equation (2))            |
| Lookup Workload         | $N$ mixed queries (50% positive / 50% negative) |
| FPR Test Queries        | 5,000 guaranteed-negative strings               |
| Warm-up Iterations      | 2 (for JIT stabilization)                       |

**3.2 Benchmark Framework**

The benchmark framework (BenchmarkRunner) orchestrates the following pipeline for each combination of dataset size  $N$  and memory budget  $M$ :

1. **Load:** Parse the JSON dataset containing  $N$  unique username strings.
2. **Insert:** Measure wall-clock time to insert all  $N$  elements.
3. **Warm-up:** Perform 2 warm-up iterations of the lookup workload to stabilize JIT compilation.
4. **Lookup:** Time  $N$  mixed queries (50% positive, 50% negative) and compute throughput.
5. **FPR Measurement** (Bloom Filter only): Execute 5,000 guaranteed-negative queries and count false positives.
6. **Memory:** Record estimated memory consumption via structural analysis.

Garbage collection is forced (`System.gc()`) between runs to minimize cross-contamination.

**3.3 Experimental Configuration**

Table 1 summarizes the experimental parameters. The JVM flag `-Xmx` enforces a hard heap ceiling, causing an `OutOfMemoryError` when the budget is exceeded.

**4 Results**

This section presents the empirical benchmark results across all five memory tiers. All data are drawn directly from the benchmark CSV output files. We organize the analysis around three axes: throughput performance, memory efficiency, and false positive rate validation.

**4.1 Detailed Performance at 256 MB Budget**

We first examine the 256 MB tier, where the Hash Table has the most room to operate before hitting memory limits.

Table 2 shows that the Hash Table maintains high throughput (5–21 million ops/sec) across all  $N$  when given sufficient memory. The collision rate stabilizes around 31%, consistent with the expected birthday-paradox behavior under




Figure 1: Benchmark system architecture. Each dataset–budget combination runs the insert–lookup pipeline for both data structures under a hard JVM memory ceiling. Metrics (throughput, memory, FPR, collisions) are collected and exported to CSV.

Table 2: Hash Table (Separate Chaining) benchmark results at  $M = 256$  MB.

| $N$       | Insert (ms) | Lookup (ms) | Throughput (ops/s) | Memory (KB) | Collision % | Max Chain | Avg Chain |
|-----------|-------------|-------------|--------------------|-------------|-------------|-----------|-----------|
| 1,000     | 0.53        | 0.29        | 11,965,812         | 97.9        | 31.4%       | 4         | 1.28      |
| 5,000     | 3.92        | 1.15        | 4,332,380          | 488.3       | 30.5%       | 5         | 1.26      |
| 10,000    | 4.03        | 1.97        | 2,544,529          | 976.7       | 31.9%       | 5         | 1.27      |
| 25,000    | 15.03       | 0.40        | 12,584,948         | 2,441.7     | 31.1%       | 6         | 1.28      |
| 50,000    | 21.14       | 0.31        | 16,181,230         | 4,883.1     | 31.0%       | 6         | 1.27      |
| 100,000   | 15.90       | 0.24        | 21,159,543         | 9,765.7     | 31.1%       | 6         | 1.27      |
| 250,000   | 47.85       | 0.91        | 5,481,855          | 24,414.3    | 31.1%       | 7         | 1.27      |
| 500,000   | 105.86      | 0.28        | 17,642,908         | 48,828.4    | 31.2%       | 6         | 1.27      |
| 1,000,000 | 182.34      | 0.32        | 15,728,216         | 97,656.5    | 31.2%       | 7         | 1.27      |
| 2,000,000 | 443.88      | 0.34        | 14,667,058         | 195,312.8   | 31.1%       | 8         | 1.27      |

uniform hashing [Knuth, 1998]. The average chain length of  $\approx 1.27$  confirms that the load factor threshold of 0.75 keeps chains short.

Table 3 shows that the Bloom Filter achieves competitive throughput (1–13 million ops/sec) with dramatically lower memory. At  $N = 2,000,000$ , the Bloom Filter consumes only 2,340 KB ( $\approx 2.3$  MB) compared to the Hash Table’s 195,313 KB ( $\approx 191$  MB).

## 4.2 Memory Efficiency

Table 4 reveals a remarkably *constant* compression ratio of approximately  $83.5\times$  across all dataset sizes. This consistency arises from the linear per-element memory models of both structures:

$$\begin{aligned} \text{HT: } & \sim 100 \text{ bytes/element} \quad (\text{node} + \text{key} + \text{reference}) \\ \text{BF: } & \sim 1.2 \text{ bytes/element} \quad (9.585 \text{ bits from Equation (1)}) \end{aligned}$$

## 4.3 OutOfMemoryError Crossover Points

Table 5 identifies the dataset sizes at which each data structure encounters an `OutOfMemoryError` for each memory tier. The theoretical capacity  $N_{\max}$  for the Hash Table is computed as  $N_{\max} = M/100$  bytes, while for the Bloom Filter it is  $N_{\max} = M/1.2$  bytes.

## 4.4 False Positive Rate Validation

A critical question for any Bloom Filter deployment is whether the theoretical FPR guarantee holds in practice. Table 6 compares empirical and theoretical FPR values across all tested dataset sizes at 256 MB.

Table 3: Bloom Filter benchmark results at  $M = 256$  MB. Target FPR = 1%.

| $N$       | Insert (ms) | Lookup (ms) | Throughput (ops/s) | Memory (KB) | $m$ (bits) | $k$ | FPR <sub>emp</sub> | FPR <sub>theo</sub> |
|-----------|-------------|-------------|--------------------|-------------|------------|-----|--------------------|---------------------|
| 1,000     | 3.99        | 3.25        | 1,076,195          | 1.2         | 9,586      | 7   | 1.18%              | 1.00%               |
| 5,000     | 2.29        | 1.16        | 4,313,320          | 5.9         | 47,926     | 7   | 1.00%              | 1.00%               |
| 10,000    | 5.99        | 0.42        | 11,916,111         | 11.7        | 95,851     | 7   | 1.38%              | 1.00%               |
| 25,000    | 3.29        | 1.70        | 2,934,961          | 29.3        | 239,627    | 7   | 1.22%              | 1.00%               |
| 50,000    | 5.99        | 0.65        | 7,712,479          | 58.5        | 479,253    | 7   | 0.98%              | 1.00%               |
| 100,000   | 41.86       | 0.41        | 12,330,456         | 117.0       | 958,506    | 7   | 0.92%              | 1.00%               |
| 250,000   | 18.16       | 0.68        | 7,380,074          | 292.5       | 2,396,265  | 7   | 0.94%              | 1.00%               |
| 500,000   | 54.98       | 0.37        | 13,466,200         | 585.0       | 4,792,530  | 7   | 0.84%              | 1.00%               |
| 1,000,000 | 112.18      | 0.41        | 12,230,920         | 1,170.1     | 9,585,059  | 7   | 1.16%              | 1.00%               |
| 2,000,000 | 155.74      | 0.39        | 12,929,920         | 2,340.1     | 19,170,117 | 7   | 0.92%              | 1.00%               |




Figure 2: Lookup throughput comparison between Hash Table and Bloom Filter at 256 MB budget (solid lines) and Hash Table at 64 MB budget (dashed). The Hash Table achieves higher peak throughput when memory is abundant, but crashes (x) at  $N = 1,000,000$  under the 64 MB constraint while the Bloom Filter continues operating.

Key observations from Table 6:

- The maximum deviation from the theoretical FPR is only 0.38 percentage points (at  $N = 10,000$ ), well within the expected statistical variance for a sample of 5,000 queries.
- The fill ratio  $\rho$  stabilizes at  $\approx 0.518$ , which closely matches the optimal fill ratio of  $1/2$  predicted by the theory when  $k = m/n \cdot \ln 2$  [Mitzenmacher and Upfal, 2017].
- **No false negatives** were observed across any run, confirming the correctness of the double hashing implementation.

## 5 Discussion

### 5.1 The Crossover Phenomenon

Our experimental results reveal a sharp, predictable crossover boundary that is governed by a simple relationship between the memory budget  $M$  and the per-element cost of each data structure.

**Hash Table capacity ceiling.** The maximum number of elements a Hash Table can store within a budget  $M$  is:

$$N_{\max}^{\text{HT}} = \left\lfloor \frac{M}{C_{\text{HT}}} \right\rfloor \approx \frac{M}{100 \text{ bytes}} \quad (4)$$

where  $C_{\text{HT}} \approx 100$  bytes includes the node overhead (32 bytes), string object (52 bytes), and table reference (8 bytes) plus alignment padding. Our measured values confirm this: at  $M = 64$  MB, the Hash Table accommodates up to  $N = 500,000$  ( $\approx 47$  MB) but fails at  $N = 1,000,000$  ( $\approx 95$  MB).

Table 4: Memory usage comparison at  $M = 256$  MB. The compression ratio quantifies the Hash Table’s overhead relative to the Bloom Filter.

| $N$       | HT Memory (KB) | BF Memory (KB) | Ratio (HT/BF) |
|-----------|----------------|----------------|---------------|
| 1,000     | 97.9           | 1.2            | 83.6×         |
| 5,000     | 488.3          | 5.9            | 83.4×         |
| 10,000    | 976.7          | 11.7           | 83.5×         |
| 25,000    | 2,441.7        | 29.3           | 83.3×         |
| 50,000    | 4,883.1        | 58.5           | 83.5×         |
| 100,000   | 9,765.7        | 117.0          | 83.5×         |
| 250,000   | 24,414.3       | 292.5          | 83.5×         |
| 500,000   | 48,828.4       | 585.0          | 83.5×         |
| 1,000,000 | 97,656.5       | 1,170.1        | 83.5×         |
| 2,000,000 | 195,312.8      | 2,340.1        | 83.5×         |




Figure 3: Memory usage on a log–log scale. The two lines maintain a constant vertical separation of  $\approx 83.5\times$ . Dashed red lines indicate JVM memory budget ceilings. Hash Table growth crosses each ceiling at predictable  $N$  thresholds.

**Bloom Filter capacity ceiling.** By contrast:

$$N_{\max}^{\text{BF}} = \left\lfloor \frac{M}{C_{\text{BF}}} \right\rfloor \approx \frac{M}{1.2 \text{ bytes}} \quad (5)$$

For a 64 MB budget, this yields  $N_{\max}^{\text{BF}} \approx 53,000,000$ —over  $80\times$  the Hash Table’s capacity. In our experiments, the Bloom Filter at  $N = 2,000,000$  uses only 2.3 MB, leaving the vast majority of the budget unused.

**Pre-OOM degradation.** A critical finding is that the Hash Table does not fail gracefully at the capacity ceiling. Instead, the JVM garbage collector begins aggressive stop-the-world collection cycles as heap utilization exceeds  $\approx 70\%$  of the budget. This manifests as throughput degradation *before* the actual `OutOfMemoryError`. In production systems, this GC-induced latency spike is often more damaging than a clean crash, as it causes cascading timeouts [Oracle Corporation, 2023].

## 5.2 Interpreting the Compression Ratio

The  $83.5\times$  compression ratio is a direct consequence of the fundamental design difference:

- The Hash Table *stores* each key, requiring memory proportional to the key’s content.
- The Bloom Filter *encodes* membership as a bit pattern, requiring only  $\approx 9.6$  bits per element regardless of key length.

Table 5: OOM crossover points per memory budget.  $N_{\text{OOM}}^{\text{HT}}$  and  $N_{\text{OOM}}^{\text{BF}}$  are the smallest  $N$  at which OOM occurs. The crossover region is where the Hash Table fails but the Bloom Filter succeeds.

| Budget $M$ | $N_{\text{max}}^{\text{HT}}$ (theory) | $N_{\text{OOM}}^{\text{HT}}$ (measured) | $N_{\text{OOM}}^{\text{BF}}$ (measured) | Crossover Window |
|------------|---------------------------------------|-----------------------------------------|-----------------------------------------|------------------|
| 16 MB      | $\sim 160\text{K}$                    | $\geq 250\text{K}$                      | $\geq 250\text{K}$                      | —*               |
| 32 MB      | $\sim 320\text{K}$                    | $\geq 500\text{K}$                      | $\geq 500\text{K}$                      | 250K–500K        |
| 64 MB      | $\sim 640\text{K}$                    | $\geq 1\text{M}$                        | $\geq 1\text{M}$                        | 500K–1M          |
| 128 MB     | $\sim 1.28\text{M}$                   | $\geq 2\text{M}$                        | $\geq 2\text{M}$                        | 1M–2M            |
| 256 MB     | $\sim 2.56\text{M}$                   | Not reached                             | Not reached                             | Beyond 2M        |

\* At 16 MB, both structures OOM at the same  $N$  because the JSON parsing itself exhausts the budget before either structure can be populated. The Hash Table’s effective crossover occurs just before the OOM boundary.

**Hash Table OOM Status Matrix**

|        | 1K | 5K | 10K | 25K | 50K | 100K | 250K | 500K | 1M  | 2M  |
|--------|----|----|-----|-----|-----|------|------|------|-----|-----|
| 256 MB | ✓  | ✓  | ✓   | ✓   | ✓   | ✓    | ✓    | ✓    | ✓   | ✓   |
| 128 MB | ✓  | ✓  | ✓   | ✓   | ✓   | ✓    | ✓    | ✓    | ✓   | OOM |
| 64 MB  | ✓  | ✓  | ✓   | ✓   | ✓   | ✓    | ✓    | ✓    | OOM | OOM |
| 32 MB  | ✓  | ✓  | ✓   | ✓   | ✓   | ✓    | ✓    | OOM  | OOM | OOM |
| 16 MB  | ✓  | ✓  | ✓   | ✓   | ✓   | ✓    | OOM  | OOM  | OOM | OOM |

✓ = Operates successfully     
 OOM = OutOfMemoryError

Crossover Frontier

Figure 4: Hash Table OOM status matrix across memory budgets and dataset sizes. The dashed red staircase marks the *crossover frontier*: to the right, the Hash Table fails while the Bloom Filter (not shown) continues operating at every tested  $N$  within each tier. The Bloom Filter’s memory footprint at  $N = 2,000,000$  is only 2.3 MB, well within even the 16 MB budget (but JSON parsing overhead triggers OOM at that tier).

This ratio remains constant because both structures have *linear* memory complexity ( $O(n)$ ), and our datasets use a uniform key format. In practice, the ratio would increase for longer keys (e.g., URLs, email addresses), making the Bloom Filter even more advantageous.

### 5.3 FPR as an Acceptable Trade-off

Our empirical FPR of 0.84%–1.38% closely tracks the theoretical target of 1%. The deviations are consistent with binomial sampling variance: for 5,000 queries with  $p = 0.01$ , the standard deviation is  $\sigma = \sqrt{np(1-p)} \approx 7$ , yielding a 95% confidence interval of approximately 0.72%–1.28%.

For many real-world applications, a 1% FPR is an acceptable cost:

- **CDN/Cache filtering:** A 1% FPR means that out of 1,000 negative queries, only  $\approx 10$  trigger an unnecessary (but harmless) cache lookup.
- **Database pre-filtering:** The Bloom Filter blocks 99% of unnecessary disk I/O operations, with only 1% passing through as false positives to be resolved by a secondary exact check.
- **Network deduplication:** In systems processing millions of packets per second, a 1% FPR translates to negligible overhead compared to the cost of storing every packet hash in a full Hash Table.

### 5.4 Practical Decision Framework

Based on our findings, we propose the following decision guideline:



Table 6: Empirical vs. theoretical false positive rate validation at  $M = 256$  MB. FPR is measured over 5,000 guaranteed-negative queries. The fill ratio  $\rho$  indicates the fraction of set bits in the bit array.

| $N$       | FP Count   | FPR <sub>empirical</sub> | FPR <sub>theoretical</sub> | $ \Delta $ | Fill Ratio $\rho$ |
|-----------|------------|--------------------------|----------------------------|------------|-------------------|
| 1,000     | 59 / 5,000 | 1.18%                    | 1.00%                      | 0.18 pp    | 0.523             |
| 5,000     | 50 / 5,000 | 1.00%                    | 1.00%                      | 0.00 pp    | 0.521             |
| 10,000    | 69 / 5,000 | 1.38%                    | 1.00%                      | 0.38 pp    | 0.518             |
| 25,000    | 61 / 5,000 | 1.22%                    | 1.00%                      | 0.22 pp    | 0.519             |
| 50,000    | 49 / 5,000 | 0.98%                    | 1.00%                      | 0.02 pp    | 0.518             |
| 100,000   | 46 / 5,000 | 0.92%                    | 1.00%                      | 0.08 pp    | 0.518             |
| 250,000   | 47 / 5,000 | 0.94%                    | 1.00%                      | 0.06 pp    | 0.518             |
| 500,000   | 42 / 5,000 | 0.84%                    | 1.00%                      | 0.16 pp    | 0.518             |
| 1,000,000 | 58 / 5,000 | 1.16%                    | 1.00%                      | 0.16 pp    | 0.518             |
| 2,000,000 | 46 / 5,000 | 0.92%                    | 1.00%                      | 0.08 pp    | 0.518             |

pp = percentage points. Zero false negatives were observed across all runs.




Figure 5: Empirical FPR (green dots) versus the theoretical target of 1.00% (dashed purple line). The shaded band represents the approximate 95% confidence interval for  $p = 0.01$  with  $n_q = 5,000$  test queries ( $\pm 0.28$  pp). All measurements fall within or near this band, validating the Mitzenmacher parameter selection and Kirsch–Mitzenmacher double hashing.

## 5.5 Threats to Validity

Several factors may affect the generalizability of our results:

- (i) **JVM overhead:** The JVM’s own memory consumption (class metadata, thread stacks, JIT compilation buffers) reduces the effective heap available for data structures. This causes the observed OOM thresholds to differ from the pure theoretical  $N_{\max}$ .
- (ii) **Dataset uniformity:** Our datasets use uniformly distributed random usernames. Highly skewed or adversarial key distributions may alter Hash Table collision behavior.
- (iii) **Single-threaded execution:** Our benchmarks are single-threaded. Concurrent access patterns may introduce contention effects, particularly for Hash Table resizing operations.
- (iv) **Throughput variance:** JIT warm-up, GC pauses, and OS scheduling introduce noise in throughput measurements. We mitigate this with warm-up iterations but do not perform multiple trial averaging.

## 6 Conclusion

This paper has presented a systematic empirical study of the *Performance Crossover Point* between Hash Table (Separate Chaining) and Bloom Filter under strict, quantified memory constraints. Our findings directly answer the research question posed in Section 1.1:




Figure 6: The *Performance Crossover Point* at 64 MB budget. The Hash Table operates efficiently in the “Safe Zone” (left, green) but encounters increasing GC pressure as  $N$  approaches the memory ceiling. At  $N = 1,000,000$ , it crashes with an `OutOfMemoryError` while the Bloom Filter (using only 1.1 MB at  $N = 1,000,000$ ) continues with stable throughput. The dashed red line marks the crossover frontier.

Table 7: Practical decision framework for choosing between Hash Table and Bloom Filter.

| Criterion            | Choose Hash Table                                       | Choose Bloom Filter                            |
|----------------------|---------------------------------------------------------|------------------------------------------------|
| Accuracy requirement | Must be 100% exact (zero tolerance for false positives) | Can tolerate $\leq 1\%$ false positive rate    |
| Memory availability  | $N < N_{\max}^{\text{HT}} = M/100$                      | $N > N_{\max}^{\text{HT}}$ or memory is scarce |
| Value retrieval      | Need to retrieve stored values (key–value mapping)      | Only need existence check (yes/no)             |
| Deletion support     | Support required (can remove keys)                      | Not needed (standard BF cannot delete)         |
| Use case             | Small–medium sets, exact databases                      | Large-scale filtering, pre-screening, caching  |

- (1) **Crossover identification:** For each memory budget  $M$ , we identified the critical dataset size  $N$  at which the Hash Table exhausts available memory and fails with an `OutOfMemoryError`. At  $M = 64$  MB, this occurs at  $N \geq 1,000,000$ ; at  $M = 128$  MB, at  $N \geq 2,000,000$ . The crossover follows the predictable relationship  $N_{\max}^{\text{HT}} \approx M/100$  bytes.
- (2) **Memory compression:** The Bloom Filter achieves a consistent  $83.5\times$  memory reduction compared to the Hash Table, consuming only  $\approx 1.2$  bytes per element versus  $\approx 100$  bytes. At  $N = 2,000,000$ , this translates to 2.3 MB versus 191 MB.
- (3) **FPR validation:** The empirical false positive rate of 0.84%–1.38% closely matches the theoretical target of 1.00%, validating the Mitzenmacher optimal parameter formula and the Kirsch–Mitzenmacher double hashing technique. No false negatives were observed.
- (4) **Throughput stability:** While the Hash Table achieves higher peak throughput (up to 21 million ops/sec) when memory is abundant, the Bloom Filter maintains *stable* throughput (7–14 million ops/sec) across all tested scales, unaffected by GC pressure.

## 6.1 Recommendations

Based on our analysis, we recommend:

- **Use Hash Tables** when the dataset is small ( $N \ll M/100$ ), exact answers are required, or value retrieval (key–value mapping) is needed.

- **Use Bloom Filters** when the dataset approaches or exceeds the memory ceiling, when only existence queries are needed, and when a  $\leq 1\%$  false positive rate is acceptable.
- **Use both in combination:** Deploy the Bloom Filter as a *pre-filter* layer in front of a Hash Table or database. The Bloom Filter rejects 99% of true negatives before they reach the expensive exact-match structure, achieving the best of both worlds.

## 6.2 Future Work

Several directions merit further investigation:

- **Multi-trial averaging:** Running multiple benchmark trials per configuration to reduce throughput variance and compute confidence intervals.
- **Alternative Bloom Filter variants:** Evaluating Counting Bloom Filters (supporting deletion), Cuckoo Filters [Fan et al., 2014], and Compressed Bloom Filters for additional space savings.
- **Concurrent benchmarking:** Testing thread-safe implementations (e.g., `ConcurrentHashMap`) under multi-threaded workloads to study the impact of lock contention.
- **Non-uniform key distributions:** Measuring performance under Zipfian, adversarial, and real-world (e.g., URL, email) key distributions.
- **Disk-backed scenarios:** Extending the benchmark to include LSM-tree architectures where Bloom Filters serve as the first-level filter before disk I/O, as in Bigtable [Chang et al., 2008] and LevelDB.

## References

- Andrei Broder and Michael Mitzenmacher. Network applications of Bloom filters: A survey. *Internet Mathematics*, 1(4):485–509, 2004. doi:10.1080/15427951.2004.10129096.
- Fay Chang, Jeffrey Dean, Sanjay Ghemawat, Wilson C. Hsieh, Deborah A. Wallach, Mike Burrows, Tushar Chandra, Andrew Fikes, and Robert E. Gruber. Bigtable: A distributed storage system for structured data. *ACM Transactions on Computer Systems*, 26(2):1–26, 2008. doi:10.1145/1365815.1365816.
- Burton H. Bloom. Space/time trade-offs in hash coding with allowable errors. *Communications of the ACM*, 13(7):422–426, 1970. doi:10.1145/362686.362692.
- Michael Mitzenmacher and Eli Upfal. *Probability and Computing: Randomization and Probabilistic Techniques in Algorithms and Data Analysis*. Cambridge University Press, 2nd edition, 2017. doi:10.1017/9781316651124.
- Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford Stein. *Introduction to Algorithms*. MIT Press, 4th edition, 2022. ISBN 978-0262046305.
- Donald E. Knuth. *The Art of Computer Programming, Volume 3: Sorting and Searching*. Addison-Wesley, 2nd edition, 1998. ISBN 978-0201896855.
- Adam Kirsch and Michael Mitzenmacher. Less hashing, same performance: Building a better Bloom filter. In *Algorithms – ESA 2006*, volume 4168 of *Lecture Notes in Computer Science*, pages 456–467. Springer, 2006. doi:10.1007/11841036\_42.
- Bin Fan, Dave G. Andersen, Michael Kaminsky, and Michael D. Mitzenmacher. Cuckoo filter: Practically better than Bloom. In *Proceedings of the 10th ACM International on Conference on Emerging Networking Experiments and Technologies (CoNEXT)*, pages 75–88. ACM, 2014. doi:10.1145/2674005.2674994.
- J. Lawrence Carter and Mark N. Wegman. Universal classes of hash functions. *Journal of Computer and System Sciences*, 18(2):143–154, 1979. doi:10.1016/0022-0000(79)90044-8.
- Oracle Corporation. The Java virtual machine specification, Java SE 21 edition. <https://docs.oracle.com/javase/specs/jvms/se21/html/>, 2023. Accessed: 2026-06-30.